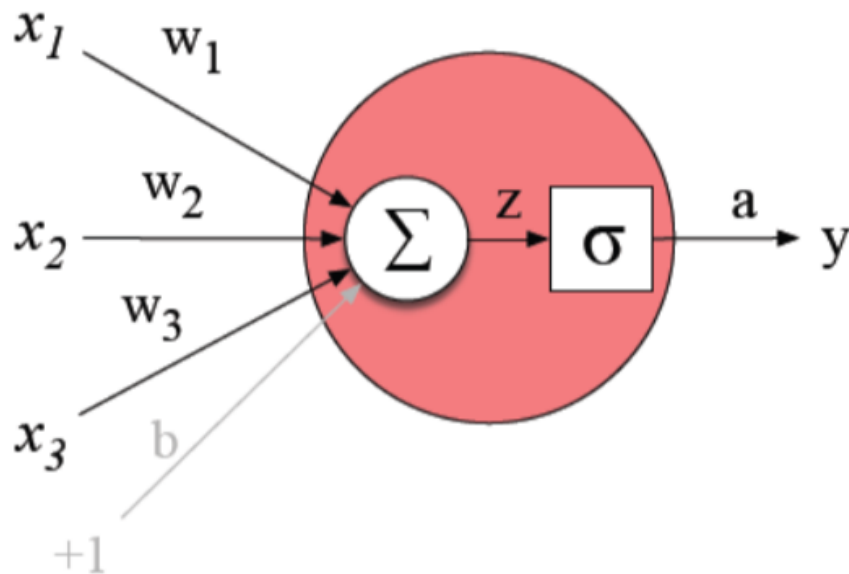


# Capitolo 5.2

## 1. Ret Neurali per nlp

- **Feed forward neural Network**

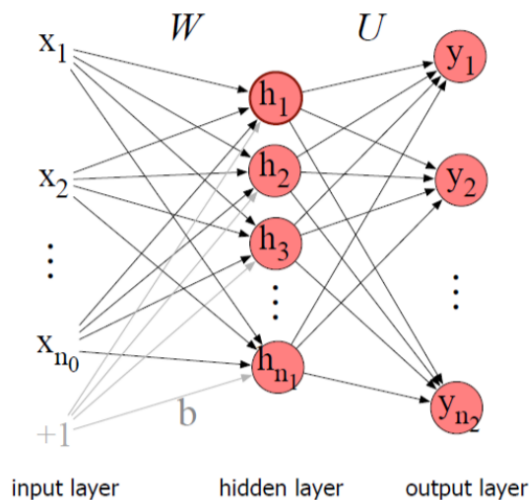
- Neural unit:



- $y = \sigma(Wx + b)$ 
  - output of the neural unit;
- $x = (x_1, \dots, x_n)$  is the input of the neural unit;
- $W$  is the weight matrix of the neural unit;
- $b$  is the bias vector of the neural unit;
- $z = Wx + b$  is the linear combination of the input and the weight matrix plus the bias vector.
- $f(z) = \sigma(z)$  is the activation function of the neural unit, where  $\sigma$  is the sigmoid function.
- other Activation function
  - tanh:  $f(z) = \tanh(z)$
  - ReLU:  $f(z) = \max(0, z)$
- the activation function is used to introduce non-linearity in the neural network.
- need non-linearity because the composition of linear functions is a linear function, and otherwise we don't have one  $W$  matrix for each layer but only one  $W$  matrix for all the layers:  $W = W_1 W_2 \dots W_n$ .
- the union of more neural unit without activation function would be the same of linear regression.

- FFNN is multilayer network with each layer composed of multiple neural units. There is at least a hidden layer, and every layer's output is the input of the next layer.

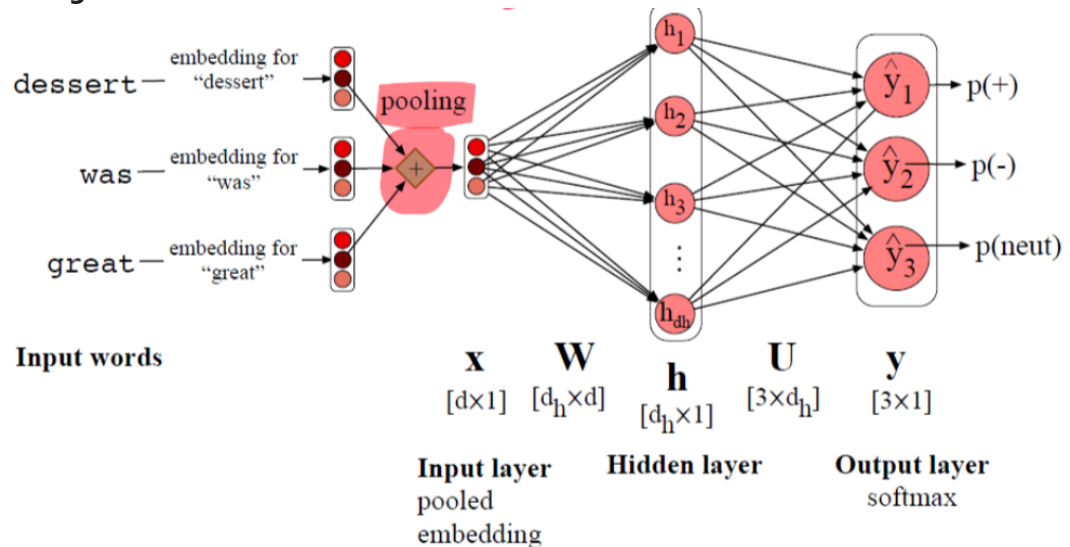
## 2-layer feed-forward network



- $h = \sigma(Wx + b)$
- $z = U \times h$
- $y = softmax(z)$
- softmax is used to normalize the output of the network in a probability distribution.
- suppose to have a network with 3 layers:
  - input layer:  $x = (x_1, x_2, x_3)$
  - hidden layer:  $h = (h_1, h_2, h_3)$ 
    - $h_1 = \sigma(W_{11}x_1 + W_{12}x_2 + W_{13}x_3 + b_1)$
    - $h_2 = \sigma(W_{21}x_1 + W_{22}x_2 + W_{23}x_3 + b_2)$
    - $h_3 = \sigma(W_{31}x_1 + W_{32}x_2 + W_{33}x_3 + b_3)$
  - output layer:  $y = (y_1, y_2, y_3)$ 
    - $y_i = \sigma(U_{i1}h_1 + U_{i2}h_2 + U_{i3}h_3 + b_i)$
  - $softmax(y) = (\frac{e^{y_1}}{\sum_{i=1}^3 e^{y_i}}, \frac{e^{y_2}}{\sum_{i=1}^3 e^{y_i}}, \frac{e^{y_3}}{\sum_{i=1}^3 e^{y_i}})$
  - the output with softmax generate a probability distribution for each class to classify.
- $W$  are the weight matrix of the layer;
- dimension of  $W$  is  $(n \times m)$ , where  $n$  is the number of neurons in the current layer and  $m$  is the number of neurons in the previous layer;
- $b$  dim is  $(n \times 1)$ ;
- $U$  is the weight matrix of the output layer;
- dim of  $U$  is  $(k \times n)$ , where  $k$  is the number of classes and  $n$  is the number of neurons in the previous layer;
- **FNN for NLP (Classification):**
  - We want to classify a sentence in a class.
  - We need to represent the sentence as a vector.

- in every case of embedding we have a matrix of size  $(V \times d)$ , where  $V$  is the size of the vocabulary and  $d$  is the dimension of the embedding.
- usually we work with sentences and not with words, so we need to represent the sentence as a vector.
- we need to represent the sentence of word as a single vector.

- **Pooling:**



- there is various way to rappresent a sequence of vectors as a single vector in order to classify the sequence.
- we can concatenate all the vectors in the sequence and use the concatenation as a single vector.
  - can be not a good idea because the dimension of the vector is too big.
- Pooling is a way to represent a sequence of vectors as a single vector.
- Given a sequence of words  $x_1, x_2, \dots, x_n$ .
- typically two type of pooling:
  - 1) **Max Pooling:**
    - $x_{pooled} = \max(x_1, x_2, \dots, x_n)$
    - $x_{pooled}$  is the vector of the max of all the words in the sequence.
  - 2) **Average pooling**
    - $x_{pooled} = \frac{1}{n} \sum_{i=1}^n e(x_i)$
    - $x_{pooled}$  is the average of all the words in the sequence.

- coming back to the classification task:

- given for ex  $x_{pooled} = \text{mean}(e(x_1), e(x_2), \dots, e(x_n))$  as the pooled representation of the sentence.
- $h = \sigma(Wx_{pooled} + b)$
- $z = U \times h$
- $y = \text{softmax}(z)$

- if we want classify a test set of  $m$  examples:

- all input sequence are packed in a matrix  $X$  of size  $(m \times d)$ , where  $m$  is the number of examples and  $d$  is the number of attribute of  $x_{pooled}$ .

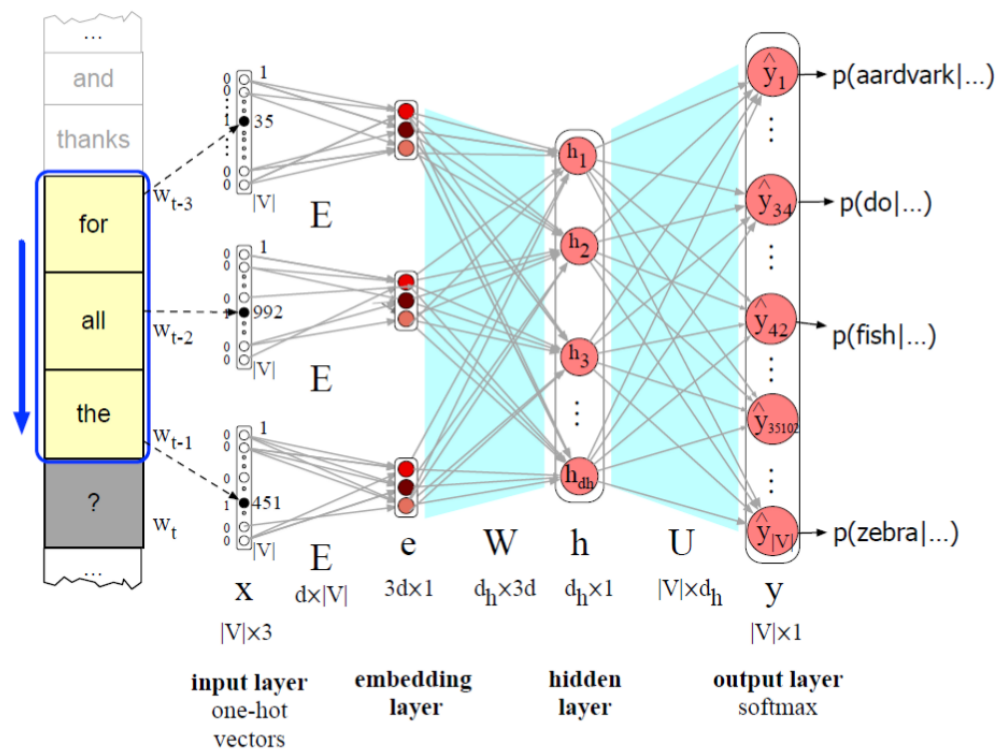
- **Feedforward Neural LM**

- the input at time  $t$  of the network is a sequence of words  $w_1, w_2, \dots, w_t$ .
- the output at time  $t$  of the network is the probability distribution of the next word  $w_{t+1}$ .
- $P(w_{t+1}|w_1, w_2, \dots, w_t)$
- assuming the words are independent:
  - $P(w_{t+1}|w_1, w_2, \dots, w_t) = \prod_{i=1}^t P(w_{i+1}|w_1, w_2, \dots, w_i) = P(w_2|w_1)P(w_3|w_1)$
- so for a sequence of word the probability is:

$$P(w_1, w_2, \dots, w_n) = \prod_{t=1}^n P(w_t|w_1, w_2, \dots, w_{t-1})$$

- Because NLM represent words by their embedding rather than word identity as in n-gram LM
- word identity in n-gram LM is a sparse representation
- sparse representation is not good for generalization because we have a lot of unseen words.
- Using embedding allows NLM to generalize better unseen words.

### Forward inference



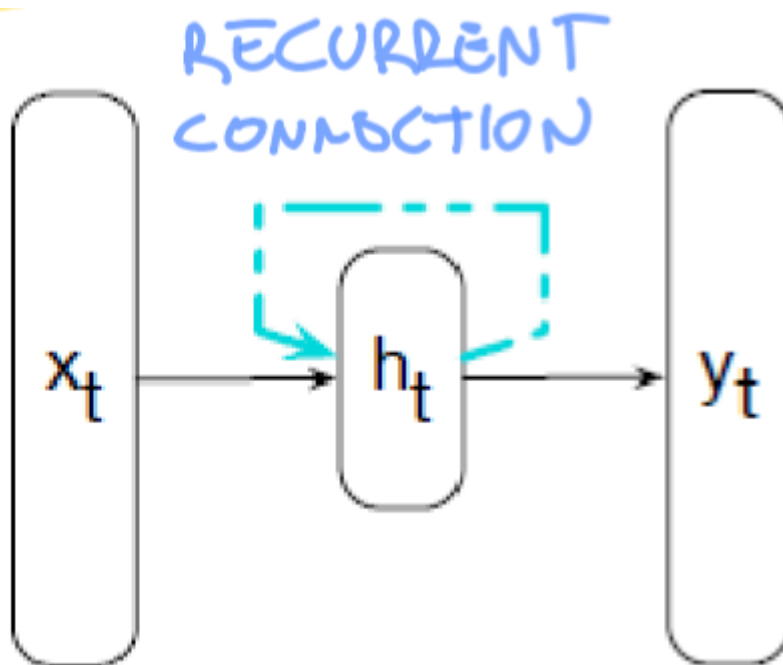
- F.inf. aka decoding:
  - is the task, given a sequence of words, to predict the next word.
  - each word is a one-hot vector of length  $V$  (one for each word in the vocabulary).
  - Embedding matrix  $E$  has a column for each word in the vocabulary and each column is the embedding of the word.
  - example of multiplication of one-hot vector with Embedding matrix:
    - $e(w_1) = E \times w_1$
    - $w_1 = [0, 0, 1]$
    - $E = \begin{bmatrix} 0.1 & 0.2 & 0.3 \\ 0.4 & 0.5 & 0.6 \\ 0.7 & 0.8 & 0.9 \end{bmatrix}$

- $e(w_1) = (0 * 0.1 + 0 * 0.4 + 1 * 0.7, 0 * 0.2 + 0 * 0.5 + 1 * 0.8, 0 * 0.3 + 0 * 0.6 + 1 * 0.9)$
- multiply one-hot vector with only one nonzero el. and embedding matrix
- this multiplication selects out relevant column of the embedding matrix for word  $w_i$ .
- Result is the embedding of the word  $w_i$ .
- **Training Neural nets:**
  - we need loss function to train the network to understand the distance between the prediction and the true label.
  - common loss function for classification task is the cross entropy loss.
  - $L = - \sum_{i=1}^k y_i \log(\hat{y}_i)$
  - where:
    - $y_i$  is the true label;
      - in case of language model is the next word of source sentence.
    - $\hat{y}_i$  is the predicted label;
      - in case of language model is the probability distribution over the vocabulary of the possible next word.
    - $k$  is the number of classes.
      - in case of language model is the size of the vocabulary.
- we need to minimize the loss function.

(Non richiesta per esame)

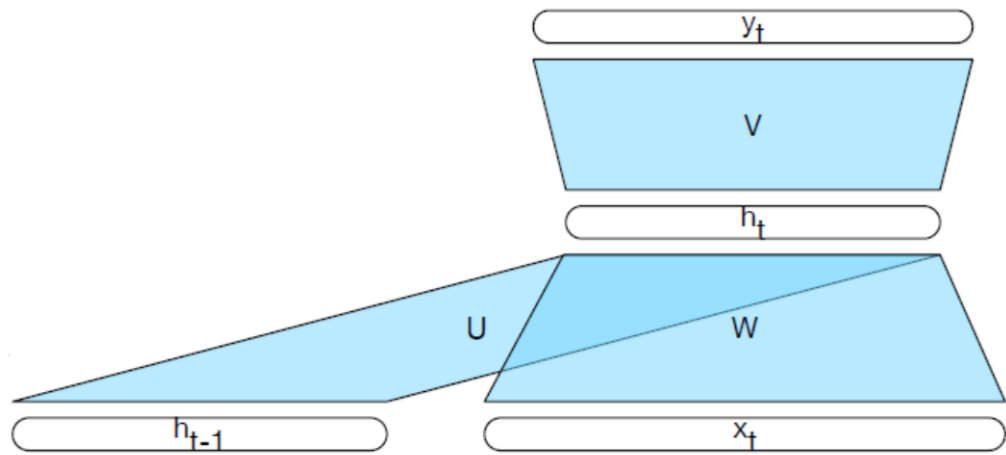
- **BACKPROPAGATION:**
  - is the algorithm to compute the gradient of the loss function with respect to the parameters of the network.
  - we need to update the parameters of the network to minimize the loss function
  - to do that we use the gradient descent.
  - we need to compute the gradient of the loss function with respect to the parameters of the network:  $\nabla_{\omega} L = (\frac{\partial L}{\partial \omega_1}, \frac{\partial L}{\partial \omega_2}, \dots, \frac{\partial L}{\partial \omega_n})$ 
    - $\omega$  is the parameter of the network.
    - we use the chain rule to compute the gradient.
    - $\frac{\partial L}{\partial \omega} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial \omega}$
    - $\omega_{new} = \omega_{old} - \eta \frac{\partial L}{\partial \omega}$
    - $\eta$  is the learning rate.
  - this process is repeated until the loss function is minimized or until the number of epochs of training is reached.

## 2. RNN



- RNN are a class of neural network that can model sequential data.
- different from feedforward neural network is the recurrent connection.
- recurrent connection allows the network to maintain information about the past.
- is a form of memory.
- memory is stored in a hidden state vector  $h$ .
- hidden state is computed as:
  - $h_t = f(Uh_{t-1} + Wx_t)$
  - $h_t$  is the hidden state vector at time  $t$ ;
  - $h_{t-1}$  is the hidden state vector at time  $t-1$ ;
  - $x_t$  is the input vector at time  $t$ ;
  - $U$  is the weight matrix for the recurrent connection;
  - $W$  is the weight matrix for the input;
  - $f$  is the activation function as sigmoid, tanh or relu.
  - the hidden state is in the hidden layer.
  - in LM the hidden state is the context of the sentence.
  - critical aspect: approach that does not impose a fixed length limit on previous context.
    - unlimited context is problematic because the hidden state vector loses information about the beginning of the sequence.
  - in fact  $h$  is a summary with loss of information of the past.
  - hidden state vector is updated at each time step;
  - so after a certain number of time steps the  $h$  will not contain information about the beginning of the sequence.

- A schematic representation of the RNN:



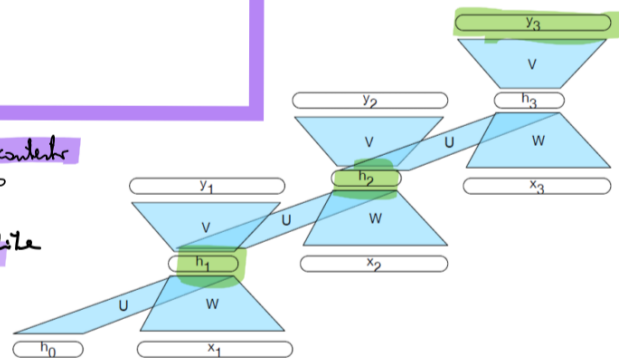
- Inference in RNN as LM

```
function FORWARDRNN(x, network) returns output sequence y
  h_0 ← 0
  for i ← 1 to LENGTH(x) do
    h_i ← g(U h_{i-1} + W x_i)
    y_i ← f(V h_i)
  return y
```

*h<sub>t</sub> è serve a tenere conto del contesto permettendo di propagare lo stato presente nella rete.*

*h è un vettore con proprietà degli input presenti e futuri*

*h<sub>3</sub> è diverso da h<sub>2</sub> perché h<sub>3</sub> è diverso da h<sub>2</sub>*

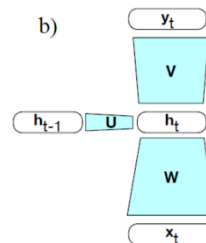
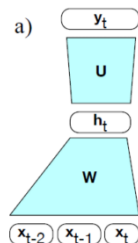


- similar to feedforward neural network.

- the difference is that the input  $x_t = [x_1, x_2, \dots, x_t]$  is a sequence of words embedded in a vector space.
- each word is a one-hot vector of length  $V$  (one for each word in the vocabulary).

feedforward  
neural language  
model

Francesco Guerra



RNN language model

*Limiting the input  
length → sequence of  
input & output  
2D array + matrix  
... forward model for  
the language model  
... the language model*

- the output  $y_t$  is a probability distribution over the vocabulary of the possible next word.
- giving input with dimension  $d_{in}$
- hidden layer with dimension  $d_h$
- output is a vector with dimension  $d_{out}$ .
- $h_t = f(Uh_{t-1} + Wx_t)$ 
  - $W$  is the weight matrix for the input of dimension  $d_h \times d_{in}$

- $U$  is the weight matrix for the recurrent connection of dimension  $d_h \times d_h$ ;
- $y_t = f(Vh_t)$ 
  - $V$  is the weight matrix for the output of dimension  $d_{out} \times d_h$ ;
  - $f$  is usually a softmax function
  - $y_t = softmax(Vh_t)$
- RNN language model process for each time step:
  - use word matrix  $E$  to retrieve the embedding of the current word  $x_t$ ;
    - $e_t = Ex_t$
  - combining with the previous hidden state  $h_{t-1}$  to compute the new hidden state  $h_t$ ;
    - $h_t = f(Uh_{t-1} + e_t)$
  - generate output layer from the hidden state  $h_t$ ;
    - $o_t = Vh_t$
  - compute the probability distribution over the vocabulary of the possible next word;
    - $y_t = softmax(o_t)$
  - Training process:
    - this is a self-supervised learning approach.
    - so the training data is unlabeled.
    - the label is the next word in the input sequence, to be compared with the output of the network.
    - compute the loss function and minimize it with backpropagation (through time).
  - using Teacher forcing: the input at time  $t$  is the true label at time  $t-1$ .
    - $x_t = y_{t-1}$
  - Weight tying:
    - Use the same embedding matrix and output weight matrix ( $V = E^T$ ).
    - $e_t = Ex_t$
    - $h_t = f(Uh_{t-1} + We_t)$
    - $y_t = softmax(E^T h_t)$
    - this is useful because:
      - $E$  and  $V$  are trained to do the same thing
      - $E$  provides a embedding for each input word
      - $V$  provide an embedding for all the next possible words
      - Using

$$V = E^T$$

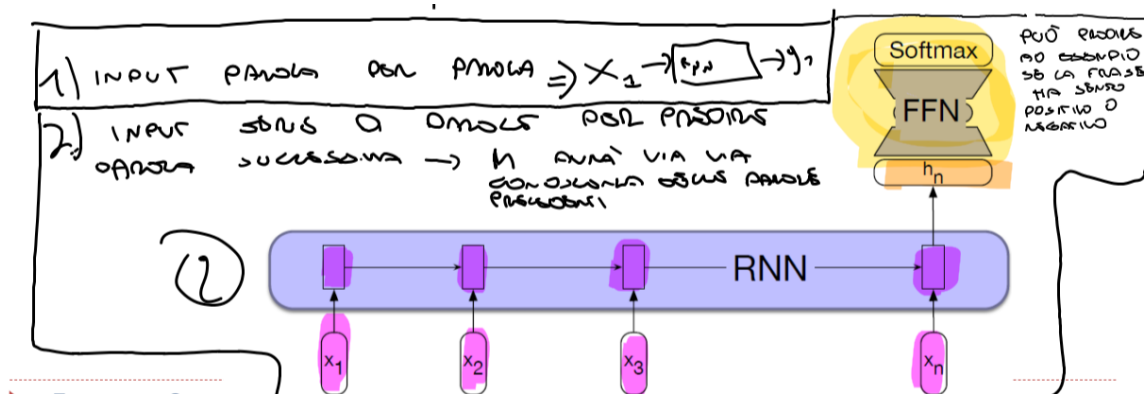
we use a single set of embedding weights for both the input and output layers.

- reduces the number of parameters to train.
- improves the performance of the model.

- RNN for Sequence classification
  - perform text classification as:



- sentiment analysis;
- spam detection;
- Give a sequence of words, predict a single label.



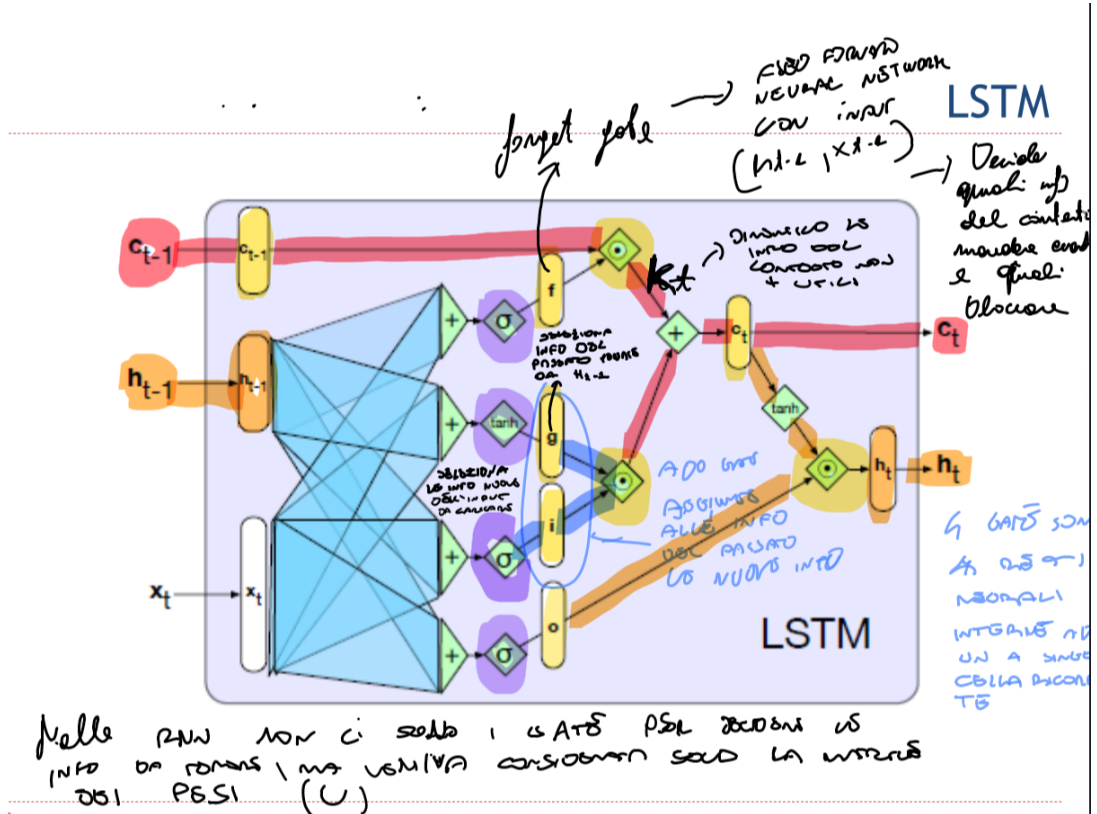
Rnn's are difficult to train because of vanishing and exploding gradient problem. Lose information about the beginning of the sequence.

### 3. LSTM:

why lstm is introduced?

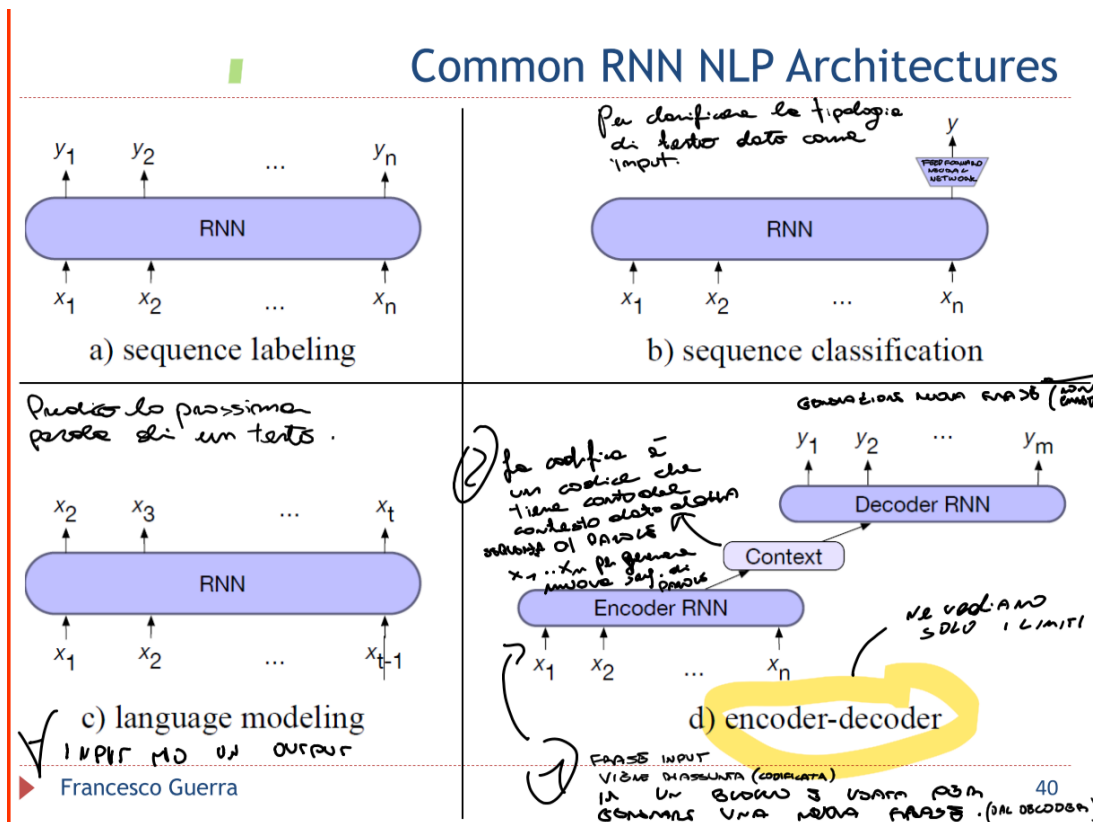
- to solve the problem of vanishing gradient and exploding gradient.
  - Maintain the context over the time.
  - removing information no longer needed from context
  - adding new information likely to be needed for later decision making.
- To manage the context, LSTM use gates.
- Gates are neural network layer that control the flow of information.
- Gates are composed of -feedforward layer;
  - sigmoid neural net layer -> push output values between 0 and 1;
  - pointwise multiplication operation
    - combining this operation with sigmoid layer -> have effect similar to a binary mask -> allow or block information flow.

- **Gates:**



- **forget gate:** compute what information to discard from the context vector.(info no longer needed)
- $f_t = \text{sigmoid}(Uf * h_t + Wf * x_t)$
- Modified context vector:  $k_t = c_{t-1} \odot f_t$  (what delete from previous context vector)
- $g_t \rightarrow$  **select candidate state:** select what new information to add from current input and previous hidden state.
  - $g_t = \tanh(Ug * h_t + Wg * x_t)$
- **input gate:**
  - $i_t = \text{sigmoid}(Ui * h_t + Wi * x_t)$
- **add gate:** select what information to add to the context vector from the candidate state.
  - $j_t = g_t \odot i_t$  (add gate)
- **output gate:** *used to decide what information is required for the current hidden state.*
  - $o_t = \text{sigmoid}(Uo * h_t + Wo * x_t)$
- **new context vector**
  - $c_t = k_t + j_t$
- **New hidden state:**
  - $h_t = o_t \odot \tanh(c_t)$

- Common RNN NLP Architectures:



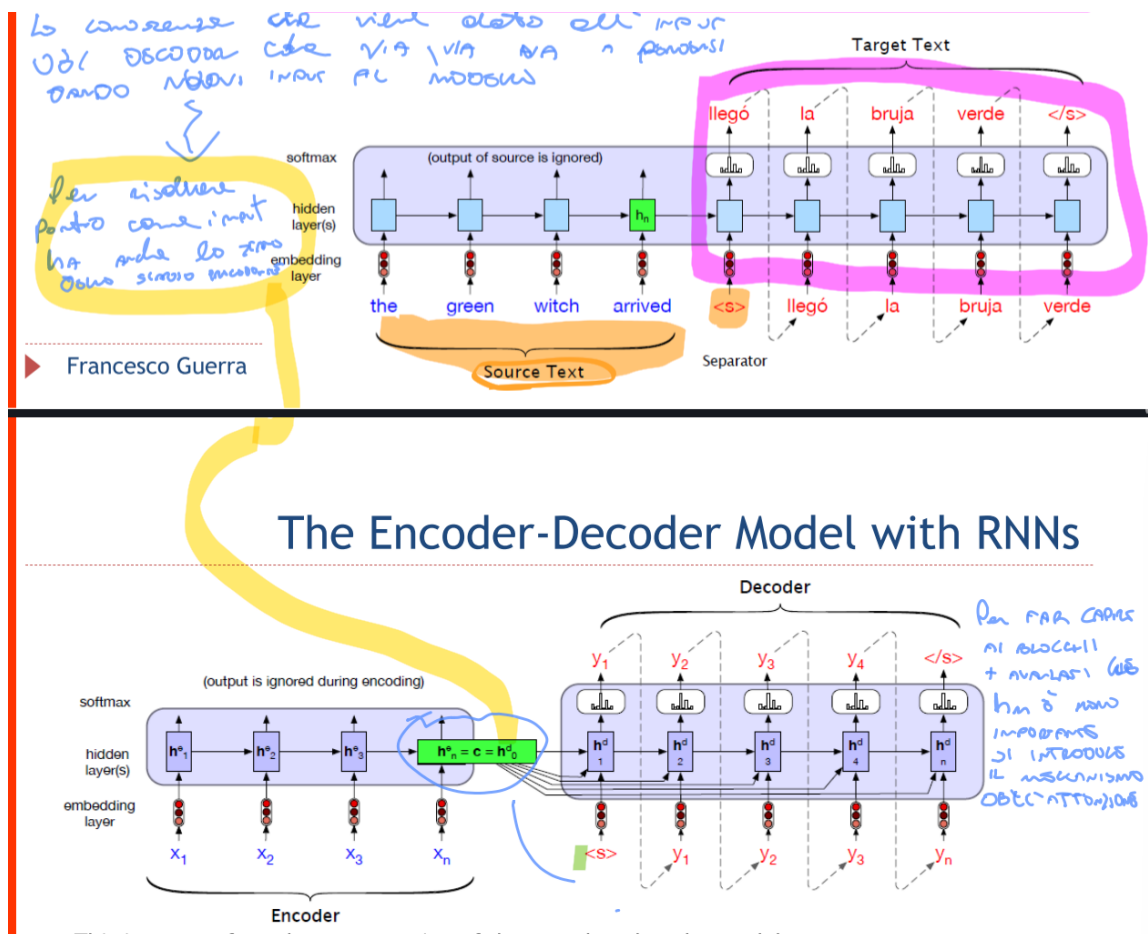
- a) Sequence labeling ( $x_1 \rightarrow y_1, x_2 \rightarrow y_2, \dots$ );
  - b) Sequence classification ( $x_1, x_2, \dots \rightarrow y$ );
  - c) language modeling ( $x_1 \rightarrow x_2, x_2 \rightarrow x_3, \dots$ );
  - d) encoder-decoder ( $x_1, x_2, \dots \rightarrow y_1, y_2, \dots$ );
- $$p(y|x) =$$
- $$p(y_1|x_1)p(y_2|x_1, y_1)p(y_3|x_1, y_2, y_3) \dots p(y_m|x_1, \dots, y_m) = p(y|x)$$
- $$= \prod_{t=1}^m p(y_t|x, y_1, \dots, y_{t-1})$$
- x is source sentence
  - y is target sentence.

- Encoder-Decoder model with RNNs:

- take input sequence arbitrary length;
- translate it into another sequence of arbitrary length.;
- task can be solved with encoder-decoder architecture:
  - summarization (input: long text, output: short text);
  - machine translation (input: text in one language, output: text in another language);
  - question answering (input: question, output: answer, as gpt3);
- Encoder (can be LSTMs, CNN, Transformer):
  - 1) take input sequence  $x = (x_1, \dots, x_n)$ ; and generate a contextualized representation of the input sequence  $h = (h_1, \dots, h_n)$ ;

- 2) context vector  $c$  is derived from the final hidden state of the encoder  $h_n$ ; -3)  
 $c$  is essence of the input to the decoder.
- Decoder (can be every model that can generate sequence):
  - 4) Decoder generates from  $c$  a arbitrary length of hidden states  
 $h = (h_1, \dots, h_m)$ ;
  - 5)  $h$  is used to generate the output sequence  $y = (y_1, \dots, y_m)$ ;
- RNN language modeling for each timestep  $t$ :
  - pass  $t - 1$  token through embedding layer;
  - Use forward inference to compute sequence of hidden states;
  - Use final hidden state to compute probability distribution over vocabulary;
    - $h_t = g(h_{t-1}, x_t)$
    - $y_t = f(h_t)$
- Thus giving:
  - $x$ : source text;
  - $y$ : target text;
- we get:

$$p(y|x) = p(y_1|x)p(y_2|x, y_1)p(y_3|x, y_1, y_2) \dots p(y_m|x, y_1, \dots, y_{m-1}) = \prod_{t=1}^m p(y_t|x, y_1, \dots$$

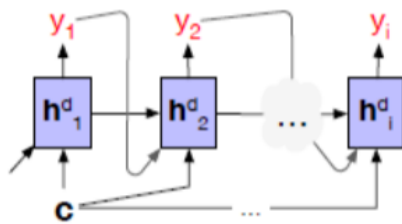


- Rnn's have problem of context influence in long term dependencies from source to target sentence.
- Introduction of attention mechanism to solve this problem.

## 4. Attention mechanism:

- allowing decoder to look at all the source words at each step of the decoding process.
- I.e. decoder get info from all hidden states of encoder, not just the last one.
  - **idea:**
    - create context vector  $c_t$  as single-fixed length, taking a weighted sum of all the hidden states of the encoder.
    - weight select relevant part of the source sentence, as the decoder generate tokens of the target sentence.
    - encoder hidden state are different for each token in decoder -> context vector is dynamically derived at each step of decoding from the hidden states of the encoder.

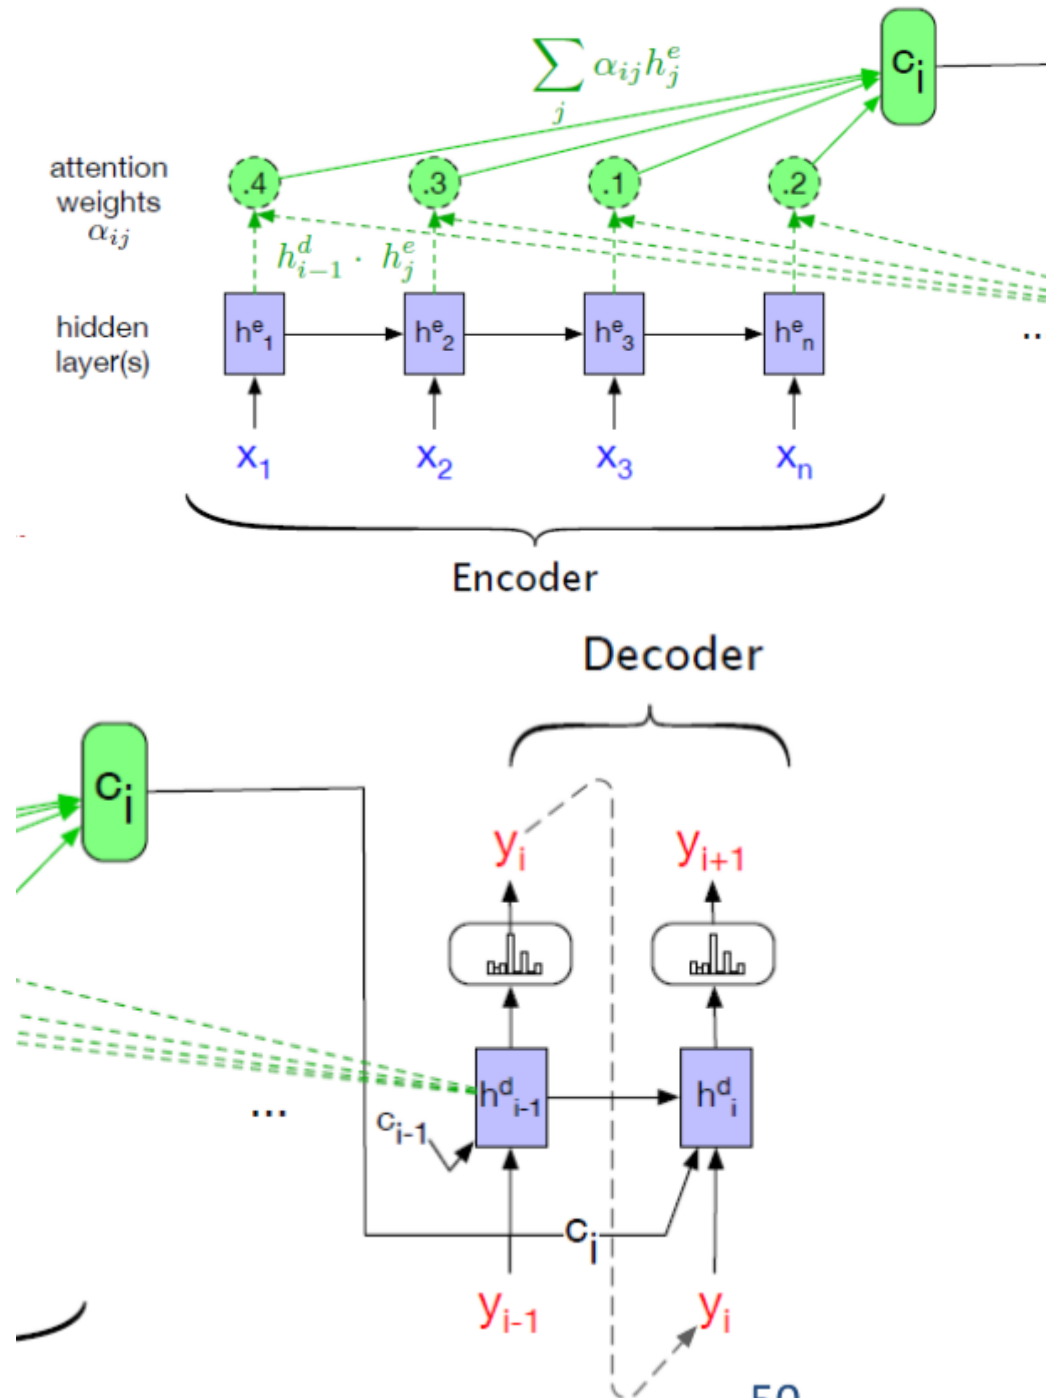
$$\mathbf{h}_t^d = g(\hat{y}_{t-1}, \mathbf{h}_{t-1}^d, \mathbf{c})$$



$$\begin{aligned} \mathbf{c} &= \mathbf{h}_n^e \\ \mathbf{h}_0^d &= \mathbf{c} \\ \mathbf{h}_t^d &= g(\hat{y}_{t-1}, \mathbf{h}_{t-1}^d, \mathbf{c}) \\ \mathbf{z}_t &= f(\mathbf{h}_t^d) \\ y_t &= \text{softmax}(\mathbf{z}_t) \end{aligned}$$

- computing  $C_i$  consider:
  - how to focus on each encoder hidden state  $h_e$ .
  - how relevant each encoder hidden state is to the current decoder hidden state  $hd_{i-1}$ .
    - introduce score to measure the relevance of each encoder hidden state during the decoding process.
- **simplest score function** (dot product attention (degree of similarity)):
  - $\text{score}(h_{e_j}, h_{d_{i-1}}) = h_{e_j} * h_{d_{i-1}}$  -> measure the similarity between j-th encoder hidden state and i-th decoder hidden state.

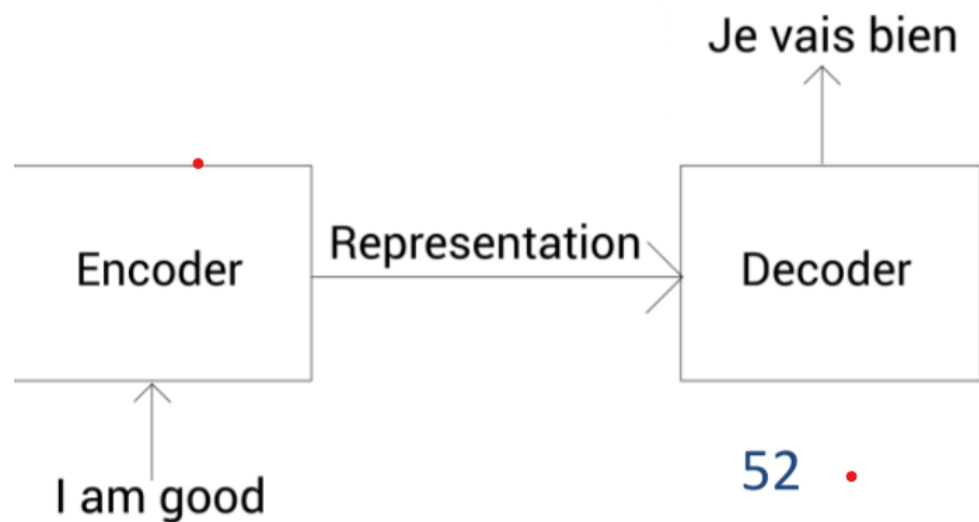
- **vector of scores:** describe how relevance each encoder hidden state is to the current decoder hidden state.
- **softmax function:**
  - $\alpha_{ij} = \text{softmax}(\text{score}(h_{e_j}, h_{d_{i-1}})) = \frac{\exp(\text{score}(h_{e_j}, h_{d_{i-1}}))}{\sum_{k=1}^n \exp(\text{score}(h_{e_k}, h_{d_{i-1}}))}$
  - give the weights to each encoder hidden state.
- **context vector:**
  - $c_i = \sum_{j=1}^n \alpha_{ij} * h_{e_j}$



## 1. Contextualized word embeddings

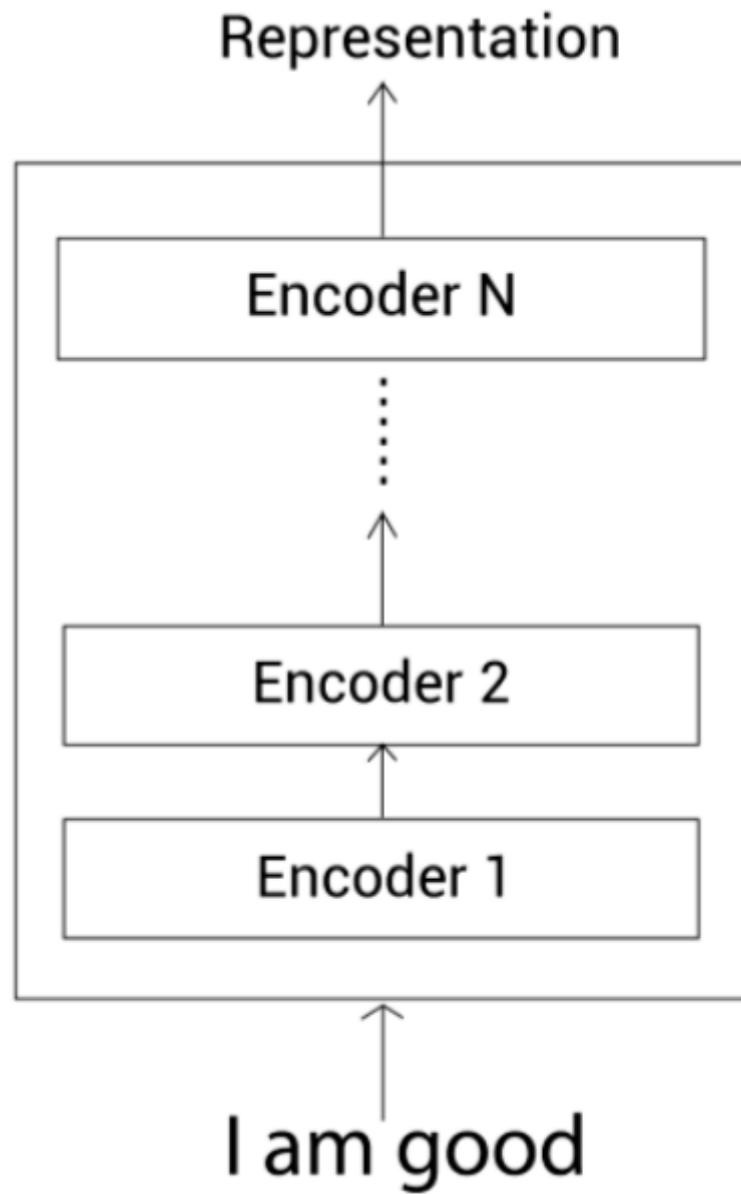
- **Transformers:**

- transformer capture long term dependencies in the input sequence by using attention mechanism.
- use self-attention to compute a representation of the input sequence.
- is encoder-decoder architecture.
  - encoder:
    - input: sequence of tokens
    - encoder learns a representation of the input sequence.
    - output: give embedding of each token in the sequence to send to decoder.
  - decoder:
    - input: embedded words from encoder
    - output: target sequence of tokens.



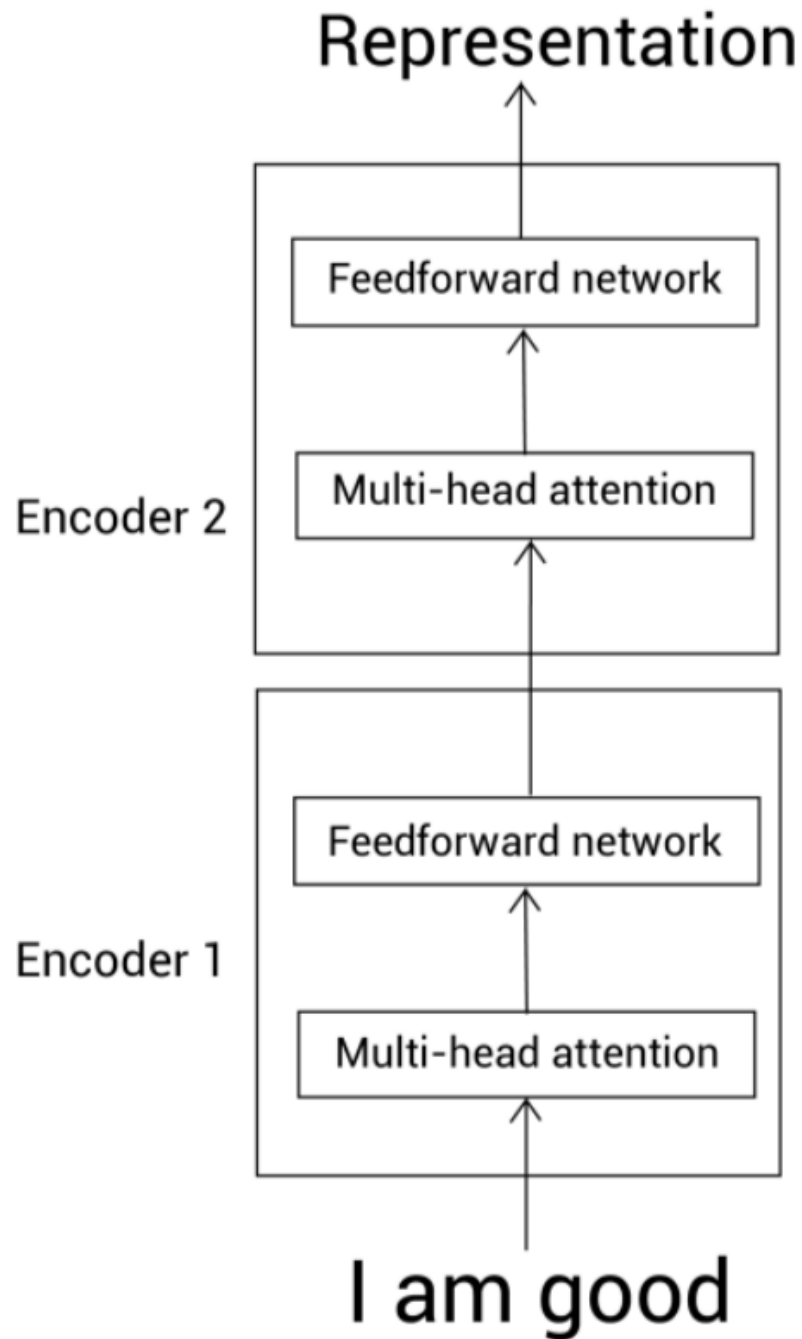
- **Trasformer encoder:**
  - consist on a stack of N encoders.

- output of each encoder is fed to the next encoder as input.



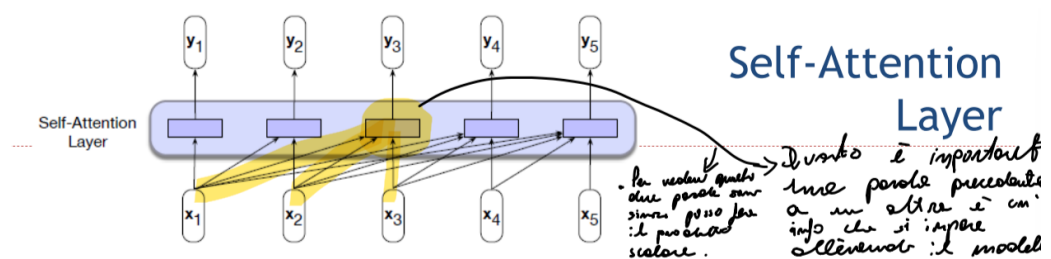
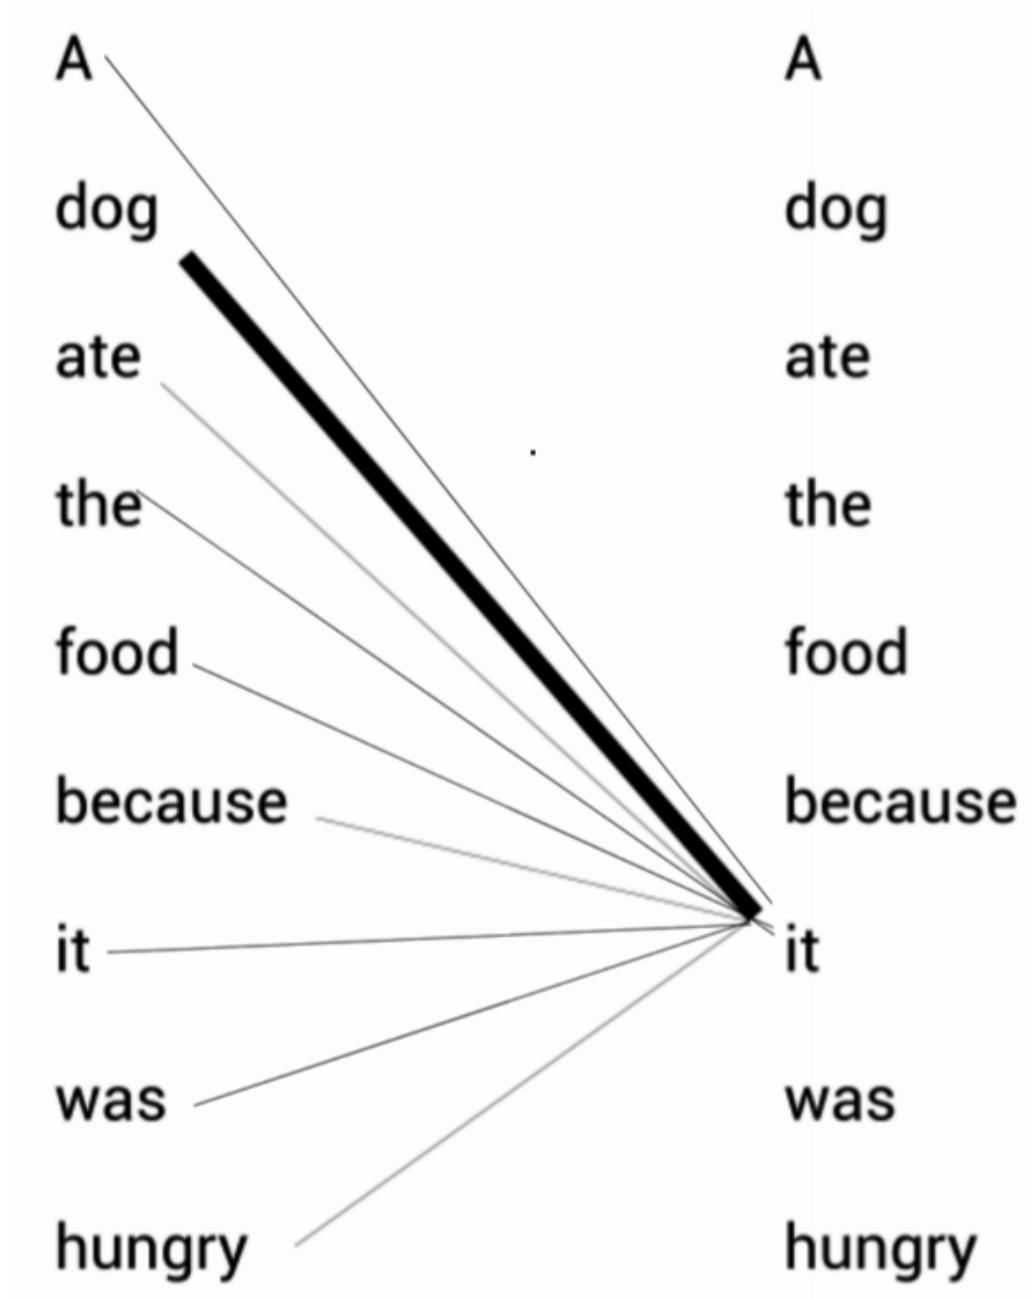
- each encoder is composed of two sublayers:
  - multi-head self-attention layer;
  - feed-forward network.





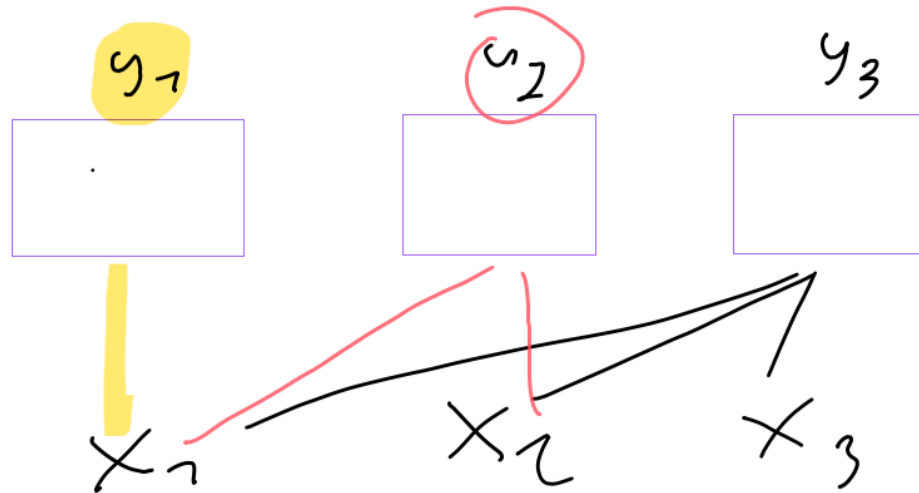
- **Self-Attention-Layer:**
  - Used to understand the relationship between different words in a sentence by computing a representation of the sentence that takes into account the

relationship between all words.



- maps input sequence of tokens  $x = (x_1, \dots, x_n)$  to sequence of vectors  $y = (y_1, \dots, y_n)$  of same length.
  - to generate output  $y_m$  the model has access to every input token  $(x_1, \dots, x_m)$ .
  - this ensure to create a LM that can be use for auto-regressive generation i.e. generate one token at a time.
- **Self-Attention-mechanism**

- provide a way to compare a word of interest to other words in the same sentence, to retrieve their relevance in the current context.
- these comparisons are used to compute an output for current input word.



- ex: to compute  $y_3$  is necessary to compare  $x_3$  with  $x_1$ ,  $x_2$  and  $x_3$  itself.
- comparison in a self attention layer with dot-product:
- $score(x_i, x_j) = x_i * x_j$
- greater the dot product, more similar the two words are.
- Score is normalized using softmax function to obtain a probability distribution over all words in the sentence.
- $\alpha_{ij} = softmax(score(x_i, x_j)) = \frac{exp(score(x_i, x_j))}{\sum_{k=1}^n exp(score(x_i, x_k))}, \forall j \leq i;$
- the probability distribution is used to compute a weighted sum of the input sequence.
- $y_i = \sum_{j \leq i} \alpha_{ij} * x_j$
- **idea:**
  - compute a representation of the input sequence by computing a weighted sum of the input sequence.
  - each input token is associated with three vectors:
    - Query vector  $q_i$ ;
    - Key vector  $k_i$ ;
    - Value vector  $v_i$ ;
  - Introducing weight matrices  $W_q$ ,  $W_k$  and  $W_v$  each of these are used to compute the query, key and value vectors for each input  $x_i$ .
  - $W_q$  have dimension  $d_{model} * d_q$ ;
  - $W_k$  have dimension  $d_{model} * d_k$ ;
  - $d_{model}$  is the dimension of the input sequence;
  - $d_q$  and  $d_k$  have to match to allow the dot product between query and key vectors;

- $W_v$  have dimension  $d_{model} * d_v$ :
- **query vector:**
  - is the current focus of attention. Is used to compute the similarity between the current input token and the previous input tokens of the sequence.
  - $q_i = W_q * x_i$ , the query vector for the i-th input token.
- **key vector:**
  - is the previous input token respect to the current input, used to compute the score.
  - $k_i = W_k * x_i$ , the key vector for the i-th input token.
- **value vector:**
  - used to compute the weighted sum of the input sequence, aka the output for the current focus of attention.
  - $v_i = W_v * x_i$ , the value vector for the i-th input token.
- **score function:**
  - $score(x_i, x_j) = \frac{q_i * k_j}{\sqrt{d_k}}$ 
    - measure the similarity between the actual focus of attention and the previous input token.
  - dot product has been scaled by  $\sqrt{d_k}$  to avoid large values of the dot product, this avoid numerical issue during training.

N

q1•k1	—∞	—∞	—∞	—∞
q2•k1	q2•k2	—∞	—∞	—∞
q3•k1	q3•k2	q3•k3	—∞	—∞
q4•k1	q4•k2	q4•k3	q4•k4	—∞
q5•k1	q5•k2	q5•k3	q5•k4	q5•k5

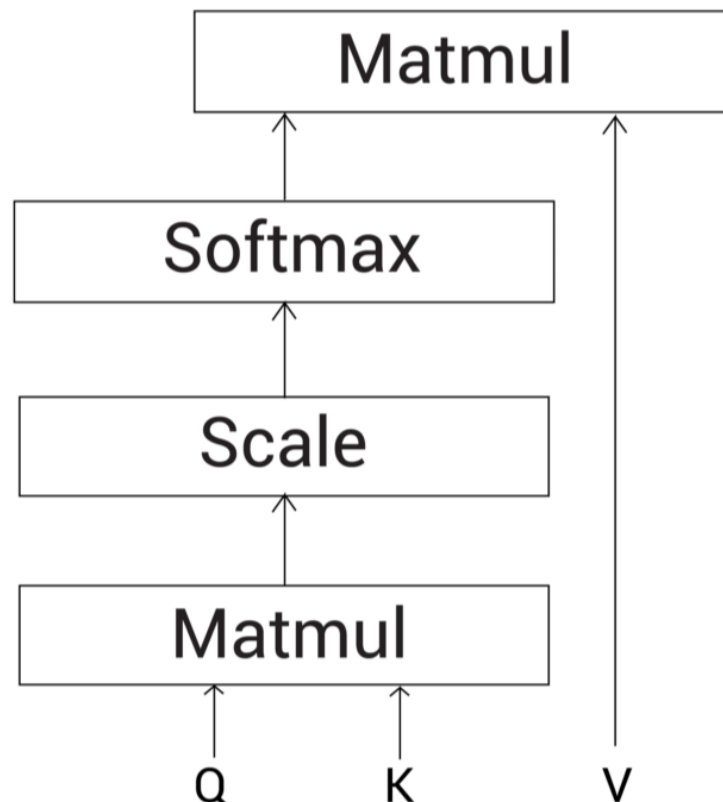
N

- Attention is quadratic in the length of the input sequence.

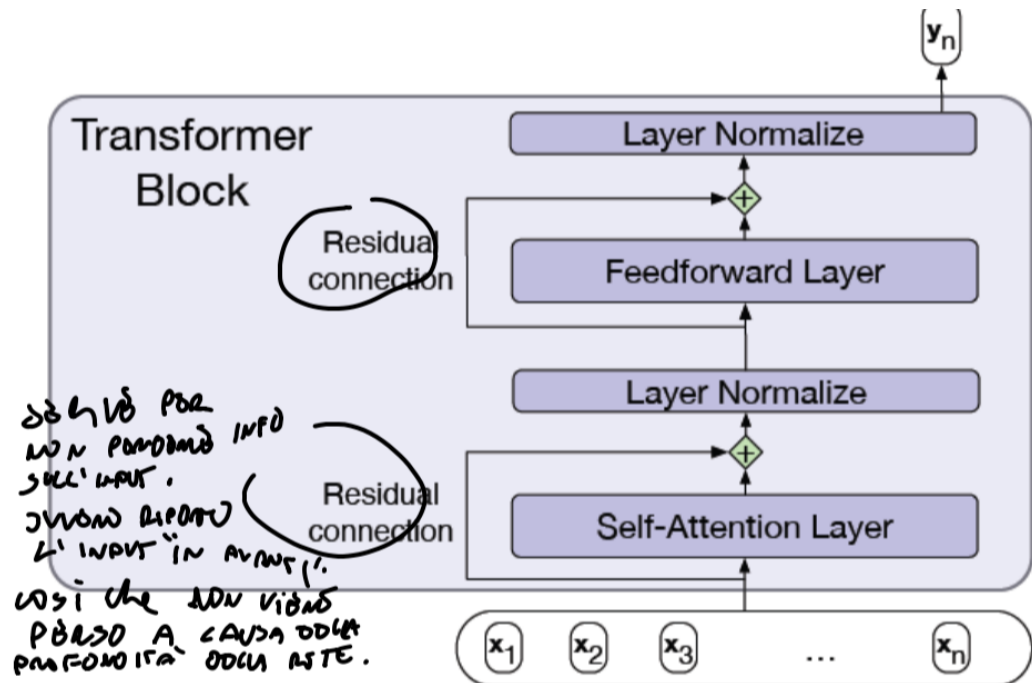
- most application have to limit the length of the input sequence to a certain number of tokens.

- **softmax function:**

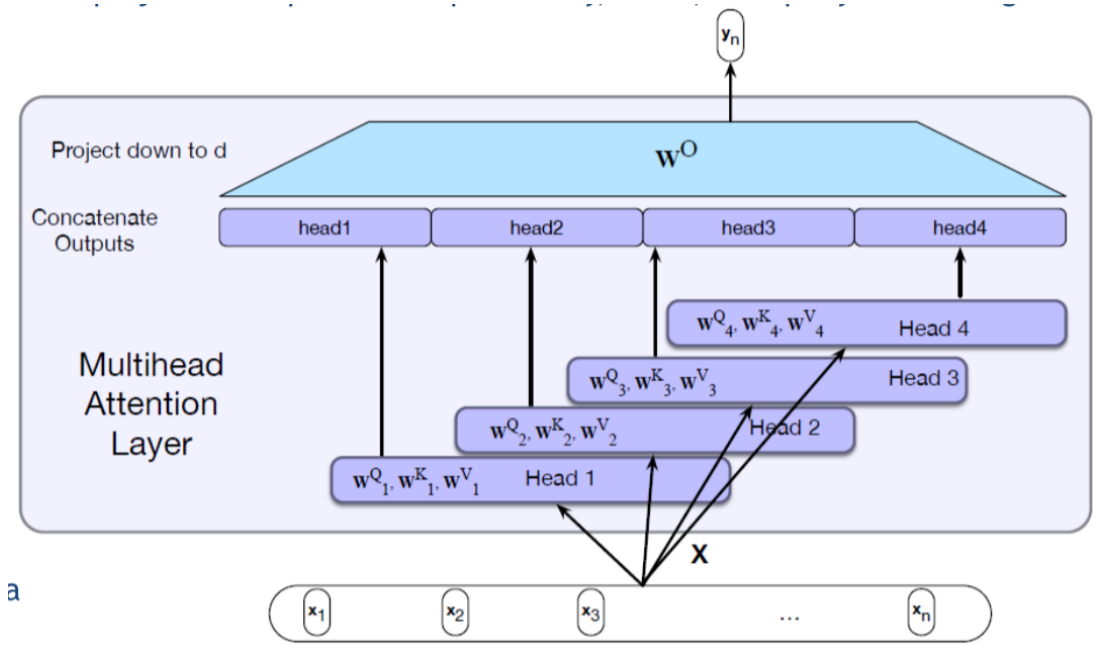
- $\text{softmax}(\text{score}(q, k)) = \frac{\exp(\text{score}(q, k))}{\sum_{j=1}^n \exp(\text{score}(q, k))}$
- give the weights to each key vector.
- **Output at time step i:**
- $y_i = \sum_{j \leq i} \text{softmax}(\text{score}(q, k)) * v$
- is the i-th output of the self attention layer.
- computed as weighted sum of the value vectors.
- each  $y_i$  is computed independently from the others.
- this allow to compute the output in parallel.
- entire process is parallelized by packing input sequence (embedding) of N tokens into a single matrix  $X \in \mathbb{R}^{N \times d_{model}}$ .
- $Q = X * W_q$
- $K = X * W_k$
- $V = X * W_v$
- $\text{SelfAttention}(Q, K, V) = \text{softmax}(\frac{Q * K^T}{\sqrt{d_k}}) * V$



- **Transformer Bloks:**



- Given input sequence  $x = (x_1, \dots, x_n)$
- pass it in self attention layer to obtain a sequence of vector  $SelfAttention(x)$
- to  $SelfAttention(x)$  is added the input sequence  $x$  with a residual connection.
- Residual connection is used to avoid to lose input information.
- after residual connection is applied a normalization layer.
- $LayerNorm(SelfAttention(x) + x)$
- $LayerNorm = \gamma * \frac{x - \mu}{\sigma} + \beta$
- $\gamma$  and  $\beta$  are learnable parameters.
- improve the training process because every input sequence will be in the same range of values.
- having input in the same range of values, the model will be able to learn faster and avoid overfitting
- **multi-head attention:**
  - Introduced to capture different relationships between different words in a sentence.
  - use multiple self-attention layers in parallel to represent the input sequence in different ways.
  - each self-attention layer is called head.
  - each head is composed of three weight matrices  $W_i^Q$ ,  $W_i^K$  and  $W_i^V$ .
  - each head is computed independently and is a different embedding of the input sequence.
  - each head is concatenated and multiplied by a weight matrix  $W^O$  to obtain the final output of the multi-head attention layer  $y$ .



- $MultiHead(Q, K, V) = Concat(head_1, \dots, head_h) * W^O$
- $head_i = Attention(Q * W_i^Q, K * W_i^K, V * W_i^V)$
- $W_i^Q \in \mathbb{R}^{d_{model} \times d_k}$
- $W_i^K \in \mathbb{R}^{d_{model} \times d_k}$
- $W_i^V \in \mathbb{R}^{d_{model} \times d_v}$
- $W^O \in \mathbb{R}^{hd_v \times d_{model}}$
- $d_k = d_v = \frac{d_{model}}{h}$
- $h$  is the number of heads.
- $d_{model}$  is the dimension of the input sequence, aka the embedding dimension for each token.
- $d_k$  and  $d_v$  are the dimension of the key and value vectors.

- **Positional Encoding:**

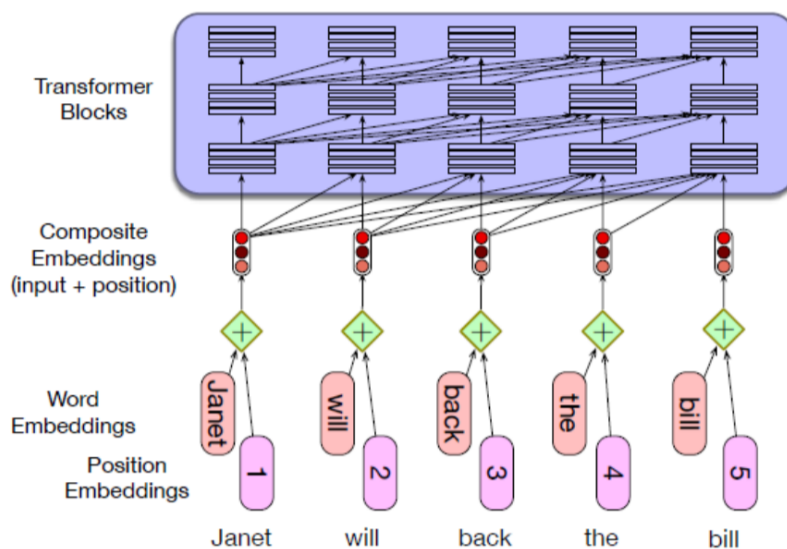
- since we fed the input sequence to the self-attention layer in parallel, the self-attention layer is not able to capture the order of the input sequence.
- to capture the order of the input sequence is added a positional encoding to the input sequence.
- positional encoding is a vector of the same dimension of the input sequence.
- each element of the positional encoding is a function of the position of the token in the input sequence.

$$\begin{aligned}
 X &= \begin{bmatrix} 1.769 & 2.22 & 3.4 & 5.8 \\ 7.3 & 9.9 & 8.5 & 7.1 \\ 9.1 & 7.1 & 0.85 & 10.1 \end{bmatrix} + \begin{bmatrix} 0 & 1 & 0 & 1 \\ 0.841 & 0.54 & 0.01 & 0.99 \\ 0.909 & -0.416 & 0.02 & 0.99 \end{bmatrix} \\
 &\quad \quad \quad X \quad \quad \quad P \\
 &= \begin{bmatrix} 1.769 & 3.22 & 3.4 & 6.8 \\ 8.14 & 10.44 & 8.51 & 8.09 \\ 10.0 & 6.68 & 0.87 & 11.09 \end{bmatrix}
 \end{aligned}$$

(non necessario per esame sapere le formule)

- $PE_{(pos,2i)} = \sin(pos/10000^{2i/d_{model}})$
- $PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{model}})$
- $d_{model}$  is the dimension of the input sequence.
- $PE_{(pos,i)}$  is the i-th element of the positional encoding for the token at position  $pos$

## Positional encoding



- **Basic idea of BERT:** (Bidirectional Encoder Representations from Transformers)
  - It is a context-based model
  - Unlike other embedding models like Word2Vec BERT is able to capture the meaning of a word based on the context in which it is used.
  - ex: Consider the following two sentences:
    - Sentence A: He got bit by Python.
    - Sentence B: Python is my favorite programming language.
    - The meaning of the word 'Python' is different in both sentences



- Context-free embedding model such as word2vec, give the same embedding of the word 'Python'.
- BERT will give different embeddings for the word 'Python' based on the context.